

Capability Tier Ladder - board one-pager

How autonomous is the agent, and what risk does that carry? Five rungs, mapped to EU AI Act risk classes. This is the shared vocabulary for any agent go/no-go decision.

The ladder

Tier	Name	Criteria	EU AI Act risk class
L0	Suggest-only	Drafts and suggestions; a human takes every consequential action.	minimal
L1	Act-with-approval	Agent proposes actions; a human approves each before execution.	limited
L2	Act-with-guardrails	Agent acts inside hard, pre-execution policy limits; humans review after the fact.	limited
L3	Act-autonomously	Agent acts unattended within a bounded domain; HITL only on flagged exceptions.	high-risk
L4	Self-directed	Agent sets and revises its own goals across domains without scoped human oversight.	prohibited

How to read it

- The risk class is the FLOOR, not the ceiling. A low-autonomy agent in a sensitive domain (hiring, credit, health, essential services) is high-risk regardless of tier.
- Promote a tier only after the evals hold at the current tier, the controls for the new tier exist, and the risk register is re-rated.
- L4 is prohibited in this handbook's stance: redesign the scope so the agent is bounded.

What each rung demands

- L0: human review of output; basic logging.
- L1: per-action approval gate; approver role + SLA; audit log of approve/reject + reason.
- L2: hard pre-execution guardrails (budget cap, allowlist, scope); post-hoc review; kill switch.
- L3: eval gate held; exception-based human oversight; immutable audit trail; on-call + incident plan.

This mapping is canonical across the handbook - the template, the diagram, chapter 12, and the README all say the same thing. A drift is a failing test.