

EU AI Act as Architecture - board one-pager

Read the EU AI Act as architecture, not paperwork. Three articles describe components you build into the system - or bolt on expensively later. Build them in.

The three articles as components

- Article 11 - Technical documentation. A living agent design spec + the capability-tier entry, generated from the system, not written about it after the fact.
- Article 12 - Record-keeping / traceability. An immutable audit log and lifecycle traces; every decision, tool call, and approval reconstructable.
- Article 14 - Human oversight. A human-in-the-loop policy and a kill switch wired into the loop - not a person told to keep an eye on it.

Risk classes

The Act classifies by use case and context: minimal, limited, high-risk, prohibited. Map each agent's capability tier to the risk class it floors at (see the capability-tier ladder one-pager). A sensitive domain raises the floor regardless of autonomy.

What it buys you

Designing to Articles 11/12/14 from day one yields an agent that is auditable, stoppable, and documented - exactly what a board wants before sign-off and what the production-readiness checklist demands. The EU lens is not overhead; it is the same discipline that makes security and evaluation trustworthy, in a form a regulator recognizes.

The differentiator

US-centric playbooks skip this. A European CPTO cannot. Treating compliance as architecture is how you ship agents in regulated markets without a last-minute scramble - and how you turn "can we even do this?" into a one-page answer.

This is an architectural aid, not legal advice. Confirm classification against current EU AI Act

guidance for your specific use case (as of June 2026 - verify before relying).